


N-Body Simulations

(Seminar: Parallelization of Physics calculations on GPUs with CUDA)

Ashish Darekar
ashish.darekar@tum.de

Shubham Khatri
shubham.khatri@tum.de

Technische Universität München
Fakultät für Physik
25th June, 2020



Overview


- Introduction: N-Body Simulations
- Particle-Particle(PP) Method - $O(N^2)$
 - Features - Force - Time Integration - Initialization
- CPU Implementation
- Improvements step I : OpenMP Parallelization
- Improvements step II : CUDA Implementation
- Performance results :
 - CPU: Serial and Serial with openMP
 - GPU: Variation in Block size - Comparison of GPU Arch. - Benchmark
- Visualisation of results - Disc of particles
- Conclusion and Improvements

Introduction : N-body simulations

“Time evolution of a system of bodies in which each body continuously interacts with every other body”

Examples:

1. Astrophysical simulation
2. Protein folding
3. Coulomb forces exerted by the atoms in a molecule
4. Turbulent fluid flow simulation




Mathematically:

$$U(\mathbf{x}_0) = \sum_i F(\mathbf{x}_0, \mathbf{x}_i)$$

References:

<https://developer.nvidia.com/gpugems/gpugems3/part-v-physics-simulation/chapter-31-fast-n-body-simulation-cuda>

Ashish Darekar & Shubham Khatri | CUDA Seminar 2020



Introduction : N-body simulations

Different Algorithms:

1. **The Particle-Particle (PP) method - $O(n^2)$**
2. The Barnes-Hut algorithm - $O(n \log n)$
3. The Particle-Mesh (PM) method - $O(n)$
4. The Particle-Particle/Particle-Mesh algorithm (P3 M) - $O(n)$
5. Fast Multipole method (FMM) - $O(n)$
6. Other methods: Hybrid Methods, Self-Consistent Field (SCF) method, Symplectic method. etc.

References:


Amara Grap's excellent web page on N-body algorithms, internet,

<http://www.amara.com/papers/nbody.html>

The Particle-Particle (PP) method - $O(n^2)$

Particle-Particle Method - Features

- Brute force technique
- Evaluate all pair-wise interactions
- $O(n^2)$ computational complexity
- $O(n^2)$ memory requirement
- Most accurate values



Simple N-Body scenario with 4 bodies


Particle-Particle Method - Force

- Force vector \mathbf{f}_{ij} on body i due to j

$$\mathbf{f}_{ij} = G \frac{m_i m_j}{\|\mathbf{r}_{ij}\|^2} \cdot \frac{\mathbf{r}_{ij}}{\|\mathbf{r}_{ij}\|},$$

- Total Force on body i

$$\mathbf{F}_i = \sum_{\substack{1 \leq j \leq N \\ j \neq i}} \mathbf{f}_{ij} = G m_i \cdot \sum_{\substack{1 \leq j \leq N \\ j \neq i}} \frac{m_j \mathbf{r}_{ij}}{\|\mathbf{r}_{ij}\|^3}.$$



Simple N-Body scenario with 4 bodies

- Addition of Softening factor:

$$\mathbf{F}_i \approx G m_i \cdot \sum_{1 \leq j \leq N} \frac{m_j \mathbf{r}_{ij}}{\left(\|\mathbf{r}_{ij}\|^2 + \epsilon^2\right)^{3/2}}.$$

Particle-Particle Method - Time Integration

- Acceleration:

$$\mathbf{a}_i \approx G \cdot \sum_{1 \leq j \leq N} \frac{m_j \mathbf{r}_{ij}}{\left(\|\mathbf{r}_{ij}\|^2 + \varepsilon^2 \right)^{3/2}}.$$


- Position and Velocity update - Euler Scheme:

$$\mathbf{v}(t + \Delta t) = \mathbf{v}(t) + \mathbf{a}(t) \Delta t$$

$$\mathbf{x}(t + \Delta t) = \mathbf{x}(t) + \mathbf{v}(t + \Delta t) \Delta t$$

Particle-Particle Method - Initialization

- Plummer model for spherical galaxy
- Simple disk galaxy : Disc of Particles
- Two galaxies: colliding disk galaxy



Particle-Particle Method - Initialization

- Mass: Heavy Central Mass

- Center Particle Position


$$P_C = \{0, 0, 0\}$$

- Disc Particle Position

$$P_{D_i} = \{R_i \cdot \cos \theta_{rand}, R_i \cdot \sin \theta_{rand}, T_i\}$$

$$T_i \in [T_{min}, T_{max}]$$

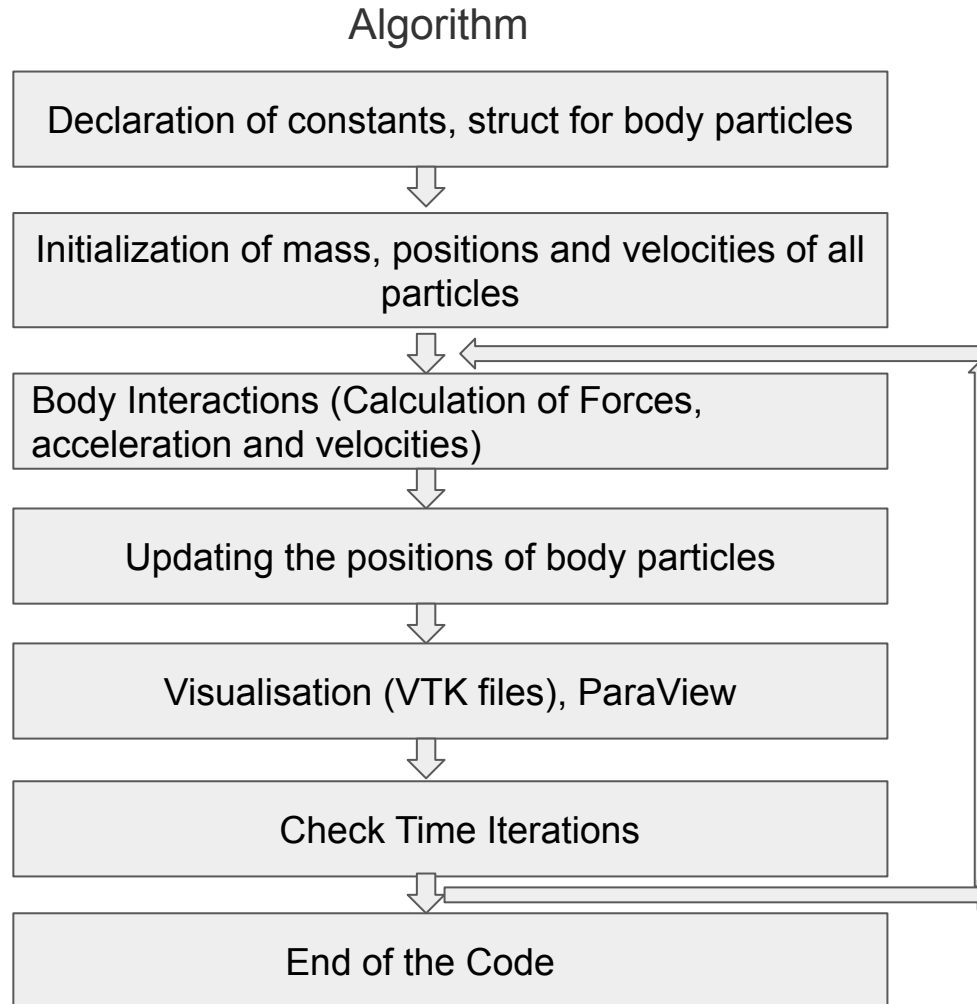
- Velocity of Particles



N-Body scenario with 4096 particles
(Disc of particles)

Serial Implementation

CPU Implementation



OpenMP Implementation

Step I: Parallelization - OpenMP

- Position update step independent.
- Position and Velocity Initialization Independent.
- Parallelize independent blocks.
- SIMD and Static schedule for performance optimization.
- Fan in reductions to increase parallelization

CUDA Implementation

Step II: Parallelization - CUDA

- Same data dependency
- Semantically same implementation
- Hypothetical Parallelism of $O(N^2)$ is possible - $N \times N$ Grid
 - Memory and Bandwidth Restriction
- Two kernels:
 - Particle Initialization
 - Position Update



Advantages:

- Initialisation on GPU - saving memcpy(HostToDevice)
- Highly parallelized initialization and computation

Step II: Parallelization - CUDA

Further consideration:




- Tiling
 - Computational Tile - $p \times p$
 - Sequential execution of tiles.
 - Better cache use



References:

<https://developer.nvidia.com/gpugems/gpugems3/part-v-physics-simulation/chapter-31-fast-n-body-simulation-cuda>

Step II: Parallelization - CUDA



References:

<https://developer.nvidia.com/gpugems/gpugems3/part-v-physics-simulation/chapter-31-fast-n-body-simulation-cuda>

Step II: Parallelization - CUDA

Further consideration:

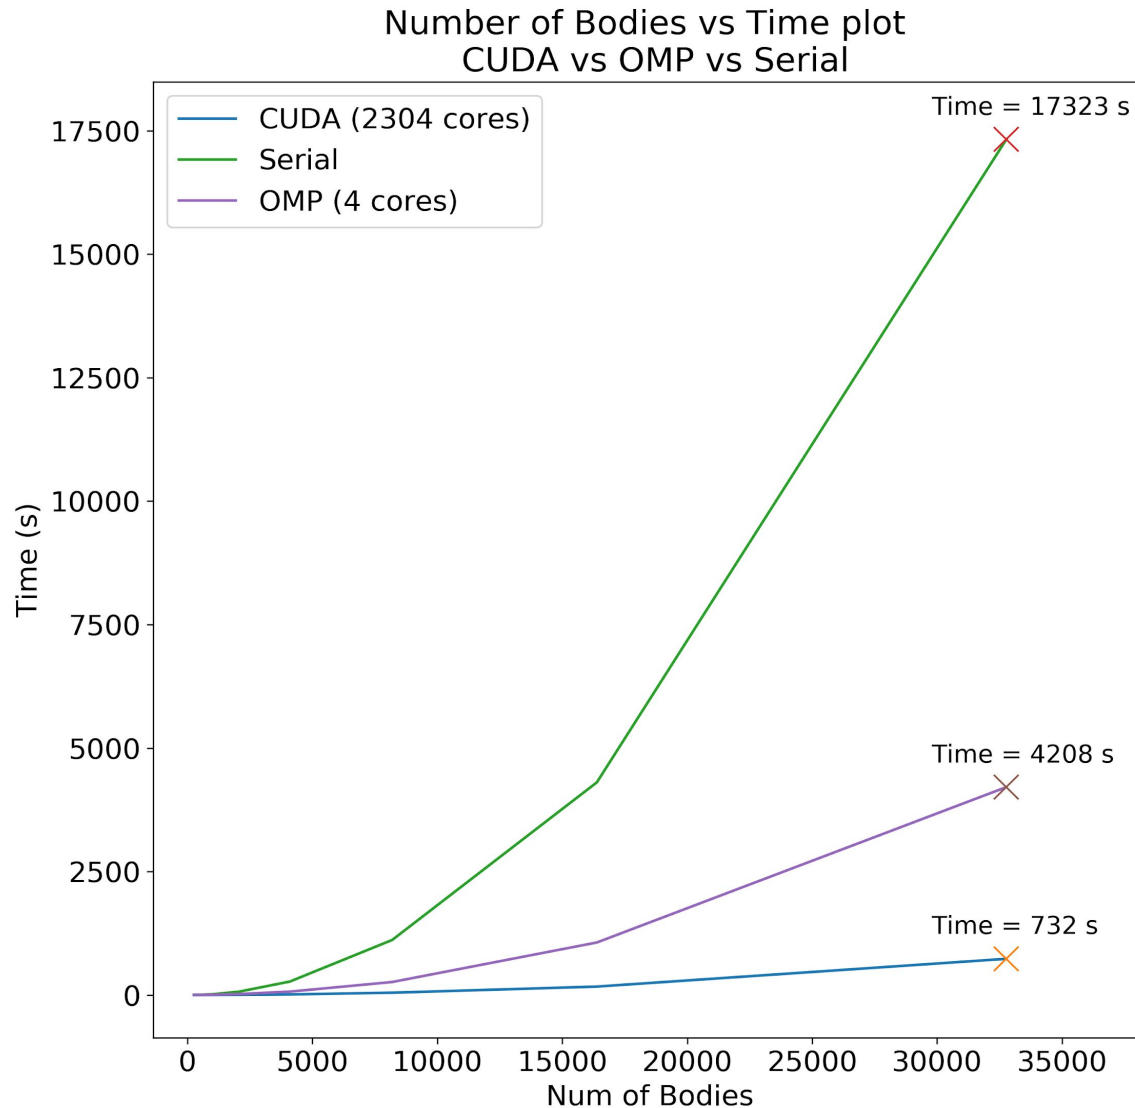
- Usage of float4
 - Coalesced memory access
 - Memory Alignment for better caching
- Loop Unrolling
 - Better thread scheduling.
 - Reduced loop control overhead.

References:

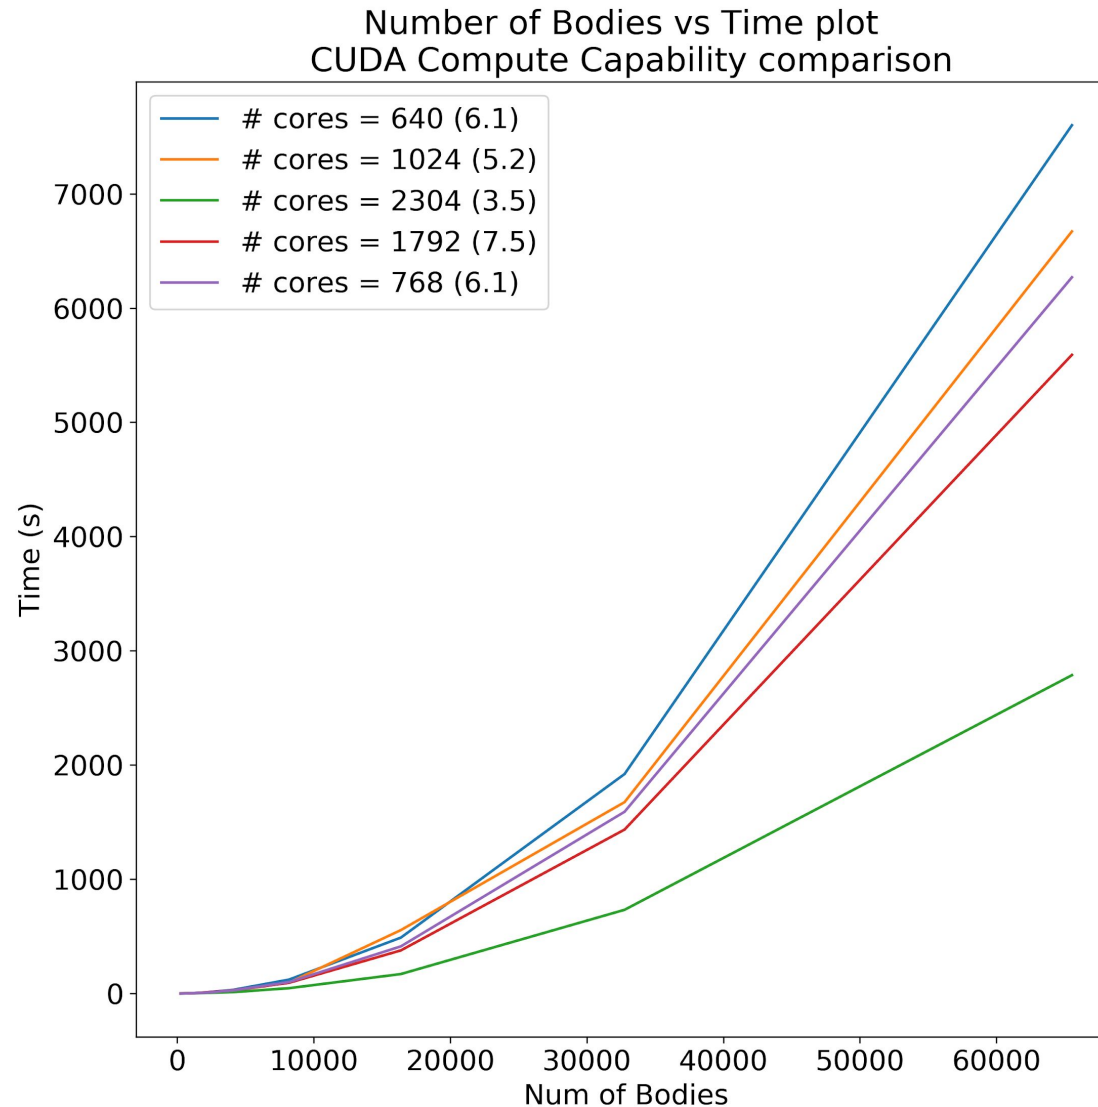
<https://developer.nvidia.com/blog/how-access-global-memory-efficiently-cuda-c-kernels/>

Performance Results:

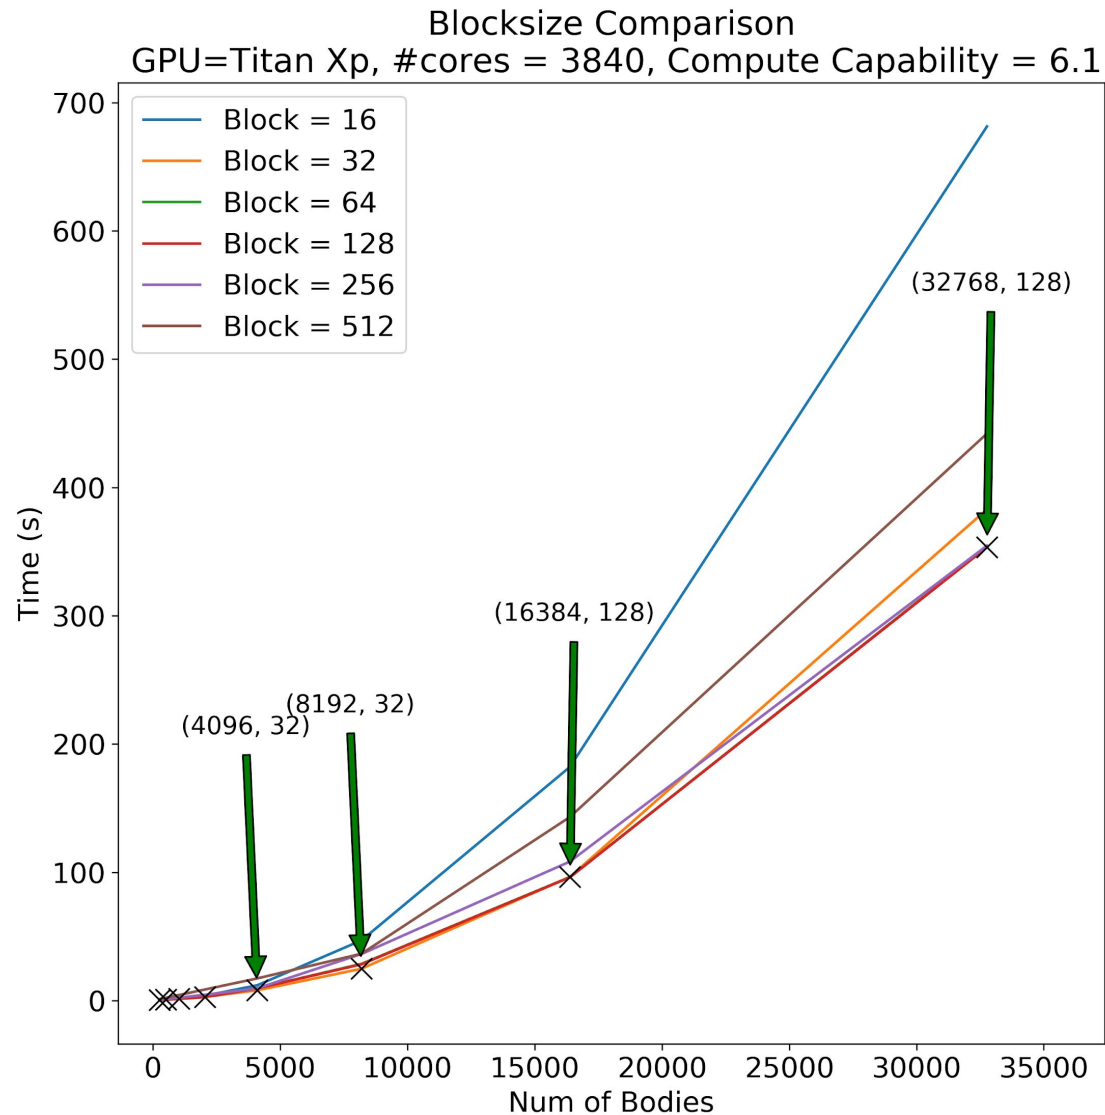
Sequential vs OpenMP vs CUDA



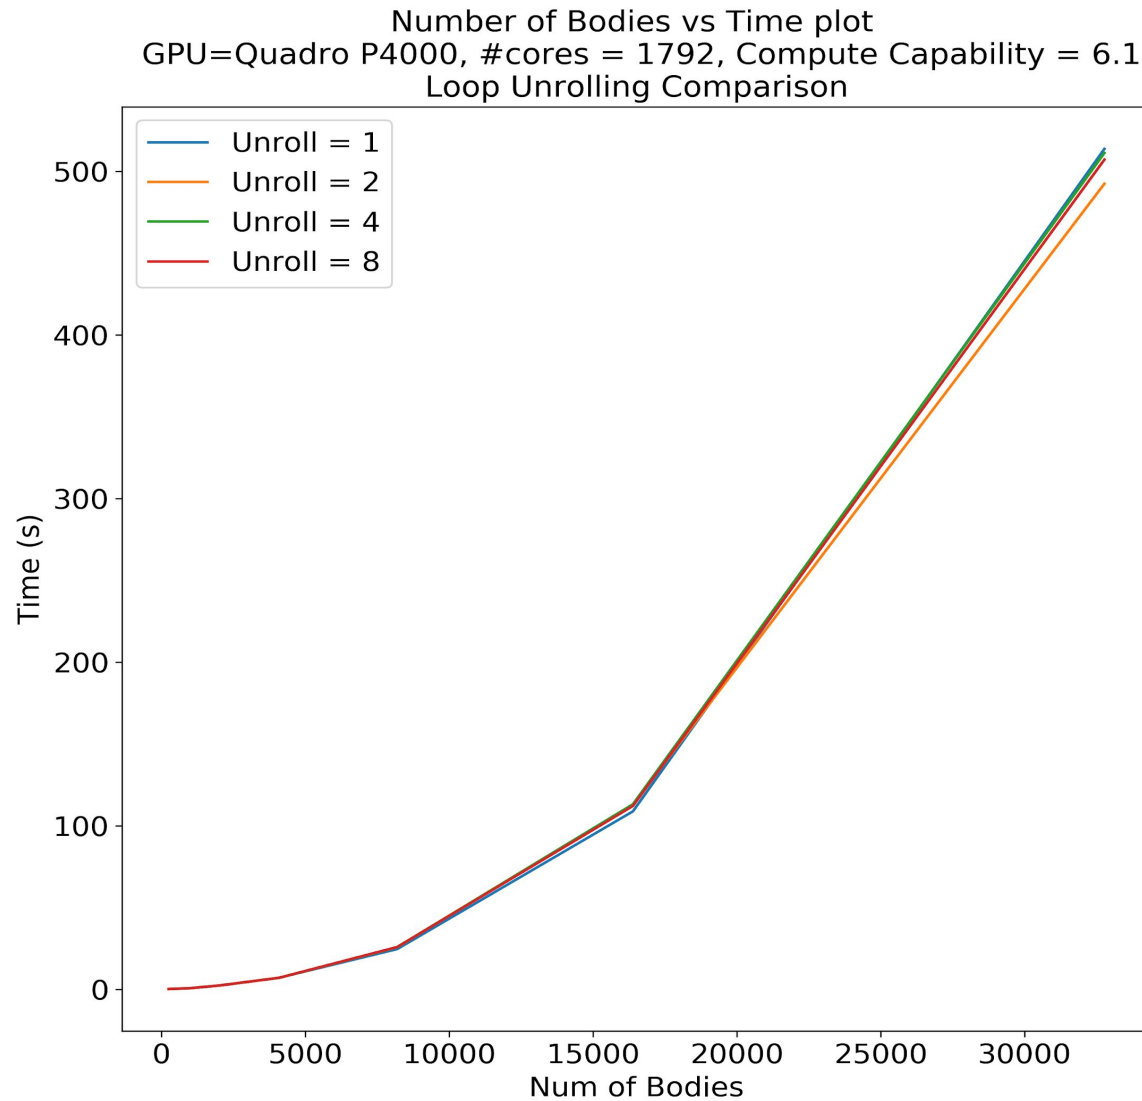
CUDA - Architecture Comparison



CUDA - Block Size Comparison



CUDA - Loop Unrolling




Benchmarking with NVIDIA Code

Hardware and parameters :


1. Machine: GeForce GTX 1050 (640 Cores)
2. Compute Capability: 6.1
3. Problem Size: 5120
4. Iterations : 10

	GFlops / s	Runtime (ms)
Our Implementation (25 Flops / interaction)	51	129.59
Nvidia (20 Flops / interaction)	1181	4.438

Visualization Results : One Particle at center



Visualization Results : Two Particles at center



Conclusion and Improvements

- PP method has computational complexity of $O(N^2)$
- Easily parallelizable using openMP, MPI or CUDA
- Performance gains through:
 - Tiling strategy (Block Size)
 - Data types - float3 and float4
 - Loop unroll
- Tree Methods like Barnes-Hut and fast Multipole Method can give better results. ($O(n \log n)$)

References

[1] CUDA N-Body simulators:

<https://developer.nvidia.com/gpugems/gpugems3/part-v-physics-simulation/chapter-31-fast-n-body-simulation-cuda>

[2] Amara Grap's excellent web page on N-body algorithms, internet,

<http://www.amara.com/papers/nbody.html>

[3] Implementation of kernel

<https://stackoverflow.com/questions/18501081/generating-random-number-within-cuda-kernel-in-a-varying-range>

<https://developer.nvidia.com/blog/easy-introduction-cuda-c-and-c/>

[4] Programming Guide CUDA - NVIDIA

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html>

<https://developer.nvidia.com/blog/how-access-global-memory-efficiently-cuda-c-kernels/>

[5] CUDA Basics

<https://www.nvidia.de/docs/IO/116711/sc11-cuda-c-basics.pdf>



Questions?